

# SEEDB: Automatically Generating Query Visualizations

Manasi Vartak  
MIT  
mvartak@mit.edu

Angela Zhang  
MIT  
angelaz@mit.edu

Aditya Parameswaran  
MIT & U. Illinois (UIUC)  
adityagp@illinois.edu

Samuel Madden  
MIT  
madden@csail.mit.edu

## ABSTRACT

Data analysts operating on large volumes of data often rely on visualizations to interpret the results of queries. However, finding the right visualization for a query is a laborious and time-consuming task. We propose SEEDB, a system that partially automates this task: given a query, SEEDB explores the space of all possible visualizations, and automatically identifies and recommends to the analyst those visualizations it finds to be most “interesting” or “useful”.

SEEDB takes advantage of current state-of-the-art research in database systems, query optimization, statistical analysis, and data visualization to simplify and guide analysts through the process of visually identifying interesting trends in the data. Little work has been done in the research community to automate this process. SEEDB aims to be a portable, time-saving, and effective tool to be included in a data analyst’s toolkit.

## 1. INTRODUCTION

Data analysts must sift through very large volumes of data to identify trends, insights, or anomalies. Given the scale of data, and the relative ease and intuitiveness of examining data visually, analysts often use visualizations as a tool to identify these trends, insights, and anomalies. However, selecting the “right” visualization often remains a laborious and time-consuming task.

We illustrate the data analysis process using an example. Consider a dataset containing sales records for a nation-wide chain of stores. Let’s say the store’s data analyst is interested in examining how the newly-introduced heating device, the “Laserwave Oven”, has been doing over the past year. The results of this analysis will inform business decisions for the chain, including marketing strategies, and the introduction of a similar “Saberwave Oven”.

The analysis workflow proceeds as follows: (1) The analyst poses a query to select the subset of data that she is interested in exploring. For instance, for the example above, she may issue the query:

```
Q = SELECT * FROM Sales WHERE Product = “Laserwave”
```

Notice that the results for this query may have (say) several million records each with several dozen attributes. Thus, directly perusing

the query result is simply infeasible. (2) Next, the analyst studies various properties of the selected data by constructing diverse views or visualizations from the data. In this particular scenario, the analyst may want to study total sales by store, quantity in stock by region, or average profits by month. To construct these views, the analyst can use operations such as binning, grouping, and aggregation, and then generate visualizations from the view. For example, to generate the view ‘total sales by store’, the analyst would group each sales record based on the store where the sale took place and sum up the sale amounts per store. This operation can easily be expressed as the familiar aggregation over group-by query:

```
Q' = SELECT store, SUM(amount) FROM Sales WHERE  
Product = “Laserwave” GROUP BY store
```

The result of the above query is a two-column table that can then be visualized as a bar-chart. Table 1 and Figure 1 respectively show an example of the results of this view and the associated visualization. To explore the query results from different perspectives, the analyst generates a large number of views (and visualizations) of the form described above. (3) The analyst then manually examines each view and decides which ones are “interesting”. This is a critical and time-consuming step. Naturally, what makes a view interesting depends on the application semantics and the trend we are comparing against. For instance, the view of Laserwave sales by store, as shown in Figure 1, may be interesting if the overall sales of all products show the *opposite* trend (e.g. Figure 2). However, the same view may be uninteresting if the sales of all products follow a similar trend (Figure 3). Thus, we posit that a view is *potentially “interesting”* if it shows a trend in the subset of data selected by the analyst (i.e., *Laserwave product-related data*) that deviates from the equivalent trend in the overall dataset. Of course, the analyst must decide if this deviation is truly an insight for this application. (4) Once the analyst has identified interesting views, the analyst may then either share these views with others, further interact with the displayed views (e.g., by drilling down or rolling up), or start afresh with a new query.

Of the four steps in the workflow described above, the ones that are especially repetitive and tedious are steps (2) and (3), where the analyst generates a large number of candidate views, and examines each of them in turn. The goal of our system, SEEDB, is to automate these labor-intensive steps of the workflow. Given a query  $Q$  indicating the subset of data that the analyst is interested in, SEEDB automatically *identifies and highlights to the analyst the most interesting views of the query results using methods based on deviation*. Specifically, SEEDB explores the space of all possible views and measures how much each view deviates from the corresponding view on the entire underlying dataset (e.g. Figure 1 vs. Figures 2 or 3.) By generating and scoring potential views automati-

This work is licensed under the Creative Commons Attribution-NonCommercial-NoDerivs 3.0 Unported License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc-nd/3.0/>. Obtain permission prior to any use beyond those covered by the license. Contact copyright holder by emailing [info@vldb.org](mailto:info@vldb.org). Articles from this volume were invited to present their results at the 40th International Conference on Very Large Data Bases, September 1st - 5th 2014, Hangzhou, China.

*Proceedings of the VLDB Endowment*, Vol. 7, No. 4

Copyright 2013 VLDB Endowment 2150-8097/13/12.

Table 1: Data: Total Sales by Store for Laser-wave

| Store             | Total Sales (\$) |
|-------------------|------------------|
| Cambridge, MA     | 180.55           |
| Seattle, WA       | 145.50           |
| New York, NY      | 122.00           |
| San Francisco, CA | 90.13            |



Figure 2: Scenario A: Total Sales by Store



Figure 1: Visualization: Total Sales by Store for Laser-wave




Figure 3: Scenario B: Total Sales by Store

ically, SEEDB effectively eliminates steps (2) and (3) that the analyst currently performs. Instead, once SEEDB recommends interesting views, the analyst can evaluate this small subset of views using domain knowledge and limit further exploration to these views.

We described our vision for SEEDB, along with the associated research challenges in a companion vision paper [?]. In this demonstration proposal, we present our first SEEDB prototype addressing some of the challenges listed in that vision paper. In particular, our current prototype of SEEDB is built as a “wrapper” that can be overlaid on any relational database system. Given any query, SEEDB leverages special optimization algorithms and the underlying DBMS to generate and recommend interesting visualizations. To do so efficiently and accurately, we must address the following challenges: (a) We must determine metrics that accurately measure the “deviation” of a view with respect to the equivalent view on the entire database (e.g., Figure 1 vs. 2), while simultaneously ensuring that SEEDB is not tied to any particular metric(s); (b) We must intelligently explore the space of candidate views. Since the number of candidate views (or visualizations) increases as the square of the number of attributes in a table (we will demonstrate this in subsequent sections), generating and evaluating all views, even for a moderately sized dataset (e.g. 1M rows, 100 attributes), can be prohibitively expensive; (c) While executing queries corresponding to different views, we must share computation as much as possible. For example, we can compute multiple views and measure their deviation all together in one query. Independent execution, on the other hand, will be expensive and wasteful; (d) Since analysis must happen in real-time, we must trade-off accuracy of visualizations or estimation of “interestingness” for reduced latency. Section 4 describes how we address these challenges.

## 2. PROBLEM STATEMENT

Given a database  $D$  and a query  $Q$ , SEEDB considers a number of views that can be generated from  $Q$  by adding relational operators. For the purposes of this discussion, we will refer to views and visualizations interchangeably, since it is straightforward to translate views into visualizations automatically. For example, there are straightforward rules that dictate how the view in Table 1 can be transformed to give a visualization like Figure 1. Furthermore, we limit the set of candidate views to those that generate a two-column result via a single-attribute grouping and aggregation (e.g. Table 1). However, SEEDB techniques can directly be used to recommend visualizations for multiple column views ( $> 2$  columns) that are generated via multi-attribute grouping and aggregation.

We consider a database  $D$  with a snowflake schema, with dimension attributes  $A$ , measure attributes  $M$ , and potential aggrega-

gate functions  $F$  over the measure attributes. We limit the class of queries  $Q$  posed over  $D$  to be those that select one or more rows from the fact table, and denote the results as  $D_Q$ .

Given such a query  $Q$ , SEEDB considers all views  $V_i$  that perform a single-attribute group-by and aggregation on  $D_Q$ . We represent  $V_i$  as a triple  $(a, m, f)$ , where  $m \in M, a \in A, f \in F$ , i.e., the view performs a group-by on  $a$  and applies the aggregation function  $f$  on a measure attribute  $m$ . We call this the *target view*.

$\text{SELECT } a, f(m) \text{ FROM } D_Q \text{ GROUP BY } a$

As discussed in the previous section, SEEDB evaluates whether a view  $V_i$  is interesting by computing the deviation between the view applied to the selected data (i.e.,  $D_Q$ ) and the view applied to the entire database. The equivalent view on the entire database  $V_i(D)$  can be expressed as shown below that we call the *comparison view*.

$\text{SELECT } a, f(m) \text{ FROM } D \text{ GROUP BY } a$

The results of both the above views are tables with two columns, namely  $a$  and  $f(m)$ . We normalize each result table into a probability distribution, such that the values of  $f(m)$  sum to 1. For our example in Table 1, the probability distribution of  $V_i(D_Q)$ , denoted as  $P[V_i(D_Q)]$ , is: (Jan: 180.55/538.18, Feb: 145.50/538.18, March: 122.00/538.18, April: 90.13/538.18). A similar probability distribution can be derived for  $P[V_i(D)]$ .

Given a view  $V_i$  and probability distributions for the target view ( $P[V_i(D_Q)]$ ) and comparison view ( $P[V_i(D)]$ ), the *utility* of  $V_i$  is defined as the distance between these two probability distributions. Formally, if  $S$  is a distance function,

$$U(V_i) = S(P[V_i(D_Q)], P[V_i(D)])$$

The utility of a view is our measure for whether the target view is “potentially interesting” as compared to the comparison view: the higher the utility, the more the deviation from the comparison view, and the more likely the view is to be interesting. Computing distance between probability distributions has been well studied, and SEEDB supports a variety of metrics to compute utility, including Earth Movers Distance, Euclidean Distance, Kullback-Leibler (K-L) Divergence, and Jensen-Shannon Distance.

Finally, we note that while other definitions of the comparison views and utility metrics are possible, for our initial exploration into visualization recommendations, we chose to focus on the intuitive definitions above.

**PROBLEM 2.1.** *Given an analyst-specified query  $Q$  on a database  $D$ , a distance function  $S$ , and a positive integer  $k$ , find  $k$  views  $V \equiv (a, m, f)$  that have the largest values of  $U(V)$  among all the views that can be represented using a triple  $(a, m, f)$ , while minimizing total computation time.*

Thus, SEEDB aims to find the  $k$  views (obtained by adding a single aggregate and group-by operator) that have the largest utility based on the function  $U$ .

### 2.1 Extensions

There are two important variations of the problem that can be addressed exactly by the same techniques discussed in the paper. These are:

- **Group-by clauses with multiple dimension attributes:** It is straightforward to extend the SEEDB techniques to group-by clauses with multiple attributes. However, for ease of exposition and visualization, we limit the number of attributes in the group-by clause to one.

- **Comparison of two-queries:** Instead of comparing the views of the input query to the entire underlying dataset, it may be more appropriate to compare them to another subset of the data (e.g. sales of “Staplers” vs. sales of “Printers”). This simply involves replacing the dataset parameter  $D$  with a second query  $Q'$ . This variation is important since it can help users find interesting differences in data. Our techniques apply to this variation unchanged and we show experimental results for this variation.

### 3. STATE-OF-THE-ART APPROACHES

Over the past few years, the research community has introduced a number of interactive data analytics tools such as ShowMe, Polaris, and Tableau [?, ?] as well as tools like Profiler allow analysts to detect anomalies in data. Unlike SEEDB, which recommends visualizations automatically, the tools place the onus on the analyst to specify the visualization to be generated. Similar visualization specification tools have also been introduced by the database community, including Fusion Tables [?] and the Devise [?] toolkit. There has been some work on browsing data cubes in OLAP, allowing analysts to find explanations, get suggestions for next cubes to visit, or identify generalizations or patterns starting from a single cube [?, ?, ?]. While we may be able to reuse the metrics from that line of work, the same techniques will not directly apply to visualizations.

#### 3.1 Data Visualization

Fusion tables [?] allows users to create visualizations layered on top of web databases; they do not consider the problem on automatically generating insightful visualizations.

Devise [?] translated user-manipulated visualizations into database queries.

Polaris/Tableau [?] is an interface for users to explore large multi-dimensional databases by visually inspecting possible visualizations. However, it places the burden of specifying visualizations to be generated onto the users. Unlike Polaris/Tableau, SEEDB automatically select visualizations of interest to the analyst.

#### 3.2 Data Cubes

There has been some work on browsing data cubes: allowing analysts to find “explanations” for why two cube values were different to various degrees of detail, to find which neighboring cubes have similar properties to the the cube under consideration, or get suggestions on what unexplored data cubes show be looked at next.

#### 3.3 Multi-Query Optimization

Multi query optimization on relationship databases allows multiple queries to be executed more efficiently by employing techniques such as executing them in parallel or executing a combined query has been studied in depth. These studies have shown that using multiple query processing algorithms are able to reduce query execution cost considerably.

#### 3.4 Utility Measurement

There are many well-established utility measurements for comparing sets of distributions. These utility measurement techniques are generally based on the distance between these distributions, since distance is an effective indication of the difference between datasets. This difference is then used to determine the utility of the comparison between these distributions. SEEDB utilizes the following distance measurements:

- **Earth Movers Distance (EMD) [?]:** Commonly used to measure differences between color histograms from images, EMD is a popular metric for comparing discrete distributions.
- **Euclidean Distance:** The L2 norm or Euclidean distance considers the two distributions to be points in a high dimensional space and measures the distance between them.
- **Kullback-Leibler Divergence (K-L divergence) [?]:** K-L divergence measures the information lost when one probability distribution is used to approximate the other one.
- **Jenson-Shannon Distance [?, ?]:** Based on the K-L divergence, this distance measures the similarity between two probability distributions.

### 4. SEEDB DESIGN

In this section, we present the SEEDB architecture, starting with an overview followed by a detailed discussion of its components. Figure 4 depicts the architecture of our system.

Our SEEDB prototype is designed as a layer on top of a traditional relational database system. While optimization opportunities are restricted by virtue of being outside the database, our design permits SEEDB to be used in conjunction with a variety of existing database systems. This allows SEEDB to be built and maintained independently of the underlying relational database system and achieve a higher level of portability.

SEEDB is comprised of two parts: a frontend and a backend. The frontend is a “thin client” that is used to issue queries and display visualizations. The backend, in contrast, performs all the computation required to generate and select views to be recommended. It is worth noting that optimization techniques are implemented on both the frontend and the backend in order to ensure a smooth user experience.

An analyst uses the frontend to issue queries to SEEDB. We provide three mechanisms for the analyst to issue queries (further discussion in Section 5). Once the analyst issues a query via the frontend, the backend takes over:

First, the Metadata Collector module queries metadata tables (a combination of database-provided and SEEDB specific tables) for information such as table sizes, column types, data distribution, and table access patterns.

The resulting metadata along with the analyst’s query is then passed to the Query Generator module. The purpose of the Query Generator is two-fold: first, it uses metadata to prune the space of candidate views to only retain the most promising ones; and second, it generates target and comparison views for each view that has not been pruned.

The SQL queries corresponding to the target and comparison views are then passed to the Optimizer module. We refer to these queries collectively as *view queries*. Next, the Optimizer module determines the best way to combine view queries intelligently so that the total execution time is minimized. (We discuss optimizations performed by SEEDB in Section 6.)

Once the Optimizer module has generated the optimized queries, SEEDB runs them on the underlying DBMS.

Results of the optimized queries are processed by the View Processor in a streaming fashion to produce results for individual views. Individual view results are then normalized and the utility of each view is computed.

Finally SEEDB selects the top  $k$  views with the highest utility and returns them to the SEEDB frontend. The frontend generates and displays visualizations for each of these view.

The analyst is able to examine the resulting visualizations, in-

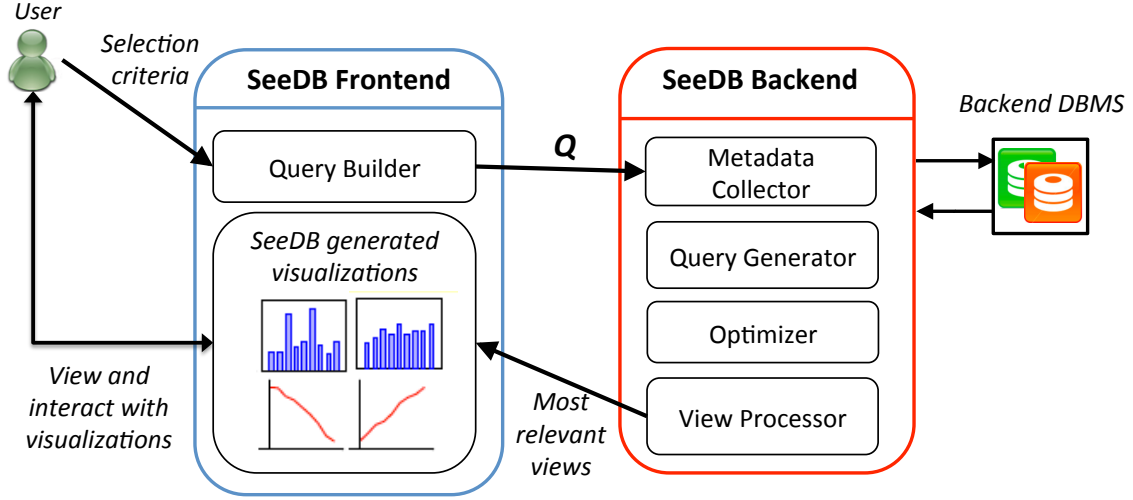


Figure 4: SeeDB Architecture

interact with SEEDB frontend to further drill into the dataset and iteratively repeat the process above to gain more insight into the dataset.

## 5. SEEDB FRONTEND

The SEEDB frontend, designed as a thin client, performs two main functions: it allows the analyst to issue a query to SEEDB, and it visualizes the results (views) produced by the SEEDB backend. To provide the analyst maximum flexibility in issuing queries, SEEDB provides the analyst with three mechanisms for specifying an input query:

- directly filling in SQL into a text box
- using a query builder tool that allows analysts unfamiliar with SQL to formulate queries through a form-based interface
- using pre-defined query templates which encode commonly performed operations, e.g., selecting outliers in a particular column.

We find that pre-defined query templates are particularly useful since analysts are often interested in anomalous data points. Addi-

tionally, SEEDB’s pre-defined templates and query building tool enable analysts who are relatively unfamiliar with the SQL syntax or the dataset being considered are still able to gain valuable insights into the dataset in one glance. This way, we have both simplified the process of insight identification as well as making it more accessible to less experienced data analysts.

Once the analyst issues a query via the SEEDB frontend, the backend evaluates various views and delivers the most interesting ones (based on utility) to the frontend. For each view delivered by the backend, the frontend creates a visualization based on parameters such as the data type (e.g. ordinal, numeric), number of distinct values, and semantics (e.g. geography vs. time series). The resulting set of visualizations is displayed to the analyst who can then easily examine these “most interesting” views at a glance, explore specific views in detail via drill-downs, and study metadata for each view (e.g. size of result, sample data, value with maximum change and other statistics). Figure 5 shows a screenshot of the SEEDB frontend (showing the query builder) in action.

After an analyst sees the resulting graphs returned by SEEDB frontend, they can also slice-and-dice views further by performing drill-downs on specific attributes in the view by (a) interactively selecting sections of the graph or (b) selecting a different group by or

Database:

donations

Table:

donations\_small

Predicates

|       | Column Name | Operator | Value            |
|-------|-------------|----------|------------------|
| where | cand_nm     | =        | Obama, Barack    |
| where | cand_nm     | in       | 'SAN FRANCISCO', |

+ Add Predicate

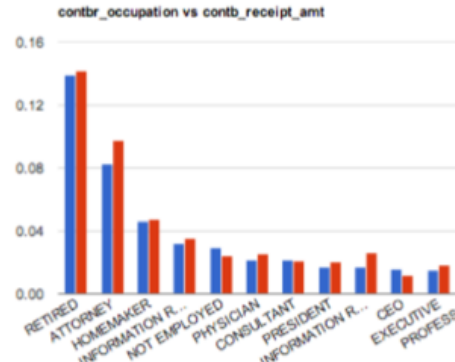




Figure 5: SeeDB Frontend: Query Builder (left) and Example Visualizations (right)

aggregate value. As such, an analyst is able to actively interact with the system to get to the most interesting results. SEEDB reduces the latency of these interactions by preemptively sending the results of all possible group-bys and aggregates to the frontend. When the analyst changes the value they want to aggregate or group by, the new graphs are rendered instantaneously using the data that has already been sent to frontend, as opposed to making another round trip to the backend and database server.

SEEDB’s frontend is designed intentionally to be simple and intuitive, such that an analyst who has limited SQL syntax knowledge and is not previously familiar with the dataset can easily identify insights in the database that they might not be able to find manually. For the advanced analysts, SEEDB allows them to directly enter queries, reduce manual labor, and speed up the insight identification process.

## 6. SEEDB BACKEND

The SEEDB backend is responsible for all the computations for generating and selecting views. As shown in Figure 4, the SEEDB backend is composed of four modules that are respectively respon-

sible for collecting metadata (Metadata Collector), pruning the view space and generating view queries (Query Generator), optimizing view queries (Optimizer), and processing query results to identify the top- $k$  interesting views (View Processor). To achieve its goal of finding the most interesting views accurately and efficiently, the SEEDB backend must not only accurately estimate the accuracy of a large number of views but also design ways in which the total processing time will be minimized. We first describe the basic SEEDB backend framework and then briefly discuss our optimizations.

### 6.1 Basic Framework

Given a user query  $Q$ , the basic approach computes all possible two-column views obtained by adding a single-attribute aggregate and group-by clause to  $Q$ . The target and comparison views corresponding to each view are then computed and each view query is executed independently on the DBMS. The query results for each view are normalized, and utility is computed as the distance between these two distributions (Section 2). Finally, the top- $k$  views with the largest utility are chosen to be displayed. The basic approach is clearly inefficient since it examines every possible view and executes each view query independently. We next discuss how

|                           |   |
|---------------------------|---|
| $T(C_i)$                  | Data type for column $C_i$                          |
| $ C_i $                   | Number of distinct values in $C_i$                  |
| $Var(C_i)$                | Variance of values in $C_i$                         |
| $Corr(C_i, C_j)$          | Correlation measure for all pairs of columns        |
| $\mathcal{H}_{i \dots k}$ | Hierarchies between columns $C_i$ to $C_k$          |
| $f_{C_i}, f_{C_i, C_j}$   | Frequency of access for each column and column pair |

Table 2: Statistics and Table Metadata

our optimizations fix these problems.

## 6.2 View Space Pruning

In practice, most views for any query  $Q$  have low utility since the target view distribution is very similar to the comparison view distribution. SEEDB uses this property to aggressively prune view queries that are unlikely to have high utility. This pruning is based on metadata about the table including data distributions and access patterns. Specifically, no expensive scans of the underlying tables are performed. In addition, we order the execution of view queries so that higher utility views can be computed before those with lower utility, thus permitting early stopping. For each table in the DBMS, we assume that statistics from Table 2 are available or can be computed cheaply. The data type for each column is numeric, categorical, ordinal, geographic, or date\_or\_time. The data type for column  $C_i$ ,  $T(C_i)$ , along with the number of distinct values  $|C_i|$  is used to determine whether the column will be treated as a *dimension* attribute or a *measure* attribute (Section ??). As before, we denote the set of dimension attributes by  $\mathcal{D}$  and measure by  $\mathcal{M}$ .

We employ the following heuristics for pruning and ordering views based on the statistics above.

- **Variance-based pruning:** Dimension attributes with low variance are likely to produce views having low utility (e.g. consider the extreme case where an attribute only takes a single value); SEEDB therefore prunes views with grouping attributes with low variance.
- **Correlated attributes:** If two dimension attributes  $a_i$  and  $a_j$  have a high degree of correlation (e.g. full name of airport and abbreviated name of airport), the views generated by grouping the table on  $a_i$  and  $a_j$  will be very similar (and have almost equal utility). We can therefore generate and evaluate a single view representing both  $a_i$  and  $a_j$ . SEEDB clusters attributes based on correlation and evaluates a representative view per cluster.
- **Bottom-up hierarchy traversal:** We observe that for a set of dimension attribute with a hierarchical structure,  $H_{C_{i \dots k}}$ , if a view  $V$  at hierarchy level  $h$  has utility  $u$ , then views at hierarchy level  $h - 1$  will have utility  $\leq u$ .
- **Access frequency-based pruning:** In tables with a large number of attributes, only a small subset of attributes are relevant to the analyst and are therefore frequently accessed for data analysis. SEEDB tracks access patterns for each table to identify the most frequently accessed columns and combinations of columns. While creating views, SEEDB uses this information to prune attributes that are rarely accessed and are thus likely to be unimportant.

**View Query Optimizations:** The second set of optimizations used by SEEDB minimizes the execution time for view queries that haven’t been pruned using the techniques described above. Since view queries tend to be very similar in structure (they differ in the aggregation attribute, grouping attribute or subset of data queried),

SEEDB uses multiple techniques to intelligently combine view queries. The ultimate goal is to minimize scans of the underlying dataset by sharing as many table scans as possible. Our strategies include:

- **Combine target and comparison view query:** Since the target view and comparison views only differ in the subset of data that the query is executed on, we can easily rewrite these two view queries as one. This simple optimization halves the time required to compute the results for a single view.
- **Combine Multiple Aggregates:** A large number of view queries have the same group-by attribute but different aggregation attributes. Therefore, SEEDB combines all view queries with the same group-by attribute into a single query. This rewriting provides a speed up linear in the number of aggregate attributes.
- **Combine Multiple Group-bys:** Since SEEDB computes a large number of group-bys, one significant optimization is to combine queries with different group-by attributes into a single query with multiple group-bys attributes. For instance, instead of executing queries for views  $(a_1, m_1, f_1)$ ,  $(a_2, m_1, f_1) \dots (a_n, m_1, f_1)$  independently, we can combine the  $n$  views into a single view represented by  $(\{a_1, a_2 \dots a_n\}, m_1, f_1)$  and post-process results at the backend. Alternatively, if the SQL GROUPING SETS functionality is available in the underlying DBMS, SEEDB can leverage that as well. While this optimization has the potential to significantly reduce query execution time, the number of views that can be combined depends on the correlation between values of grouping attributes and system parameters like the working memory. Given a set of candidate views, we model the problem of finding the optimal combinations of views as a variant of bin-packing and apply ILP techniques to obtain the best solution.
- **Sampling:** For datasets of large size, an optimization that affects performance significantly is employing sampling: we construct a sample of the dataset that can fit in memory and run all view queries against the sample. However, as expected, the sampling technique and size of the sample both affect view accuracy.
- **Parallel Query Execution:** The final optimization that SEEDB employs is taking advantage of parallel query execution at the DBMS to reduce total latency. We observe that as the number of queries executed in parallel increases, the total latency decreases at the cost of increased per query execution time.

It is possible to collect the above statistics at the dataset level too, as opposed to the entire table level. The advantage of table level statistics is that they have to be computed only once per table; however, dataset-level statistics are more accurate since they only consider the specific parts of the table.

## 7. PROPOSED USER STUDY

We propose to measure the effectiveness of SEEDB through hands-on interaction with a variety of datasets. Our goals are two fold: (1) determine the utility of SEEDB in surfacing interesting trends for a query and (2) determine whether we can return high quality views efficiently for a range of datasets. We will use four different

datasets in our user study:

- **Store Orders dataset** [?]: This dataset is often used by Tableau [?] as a canonical dataset for business intelligence applications. It consists of information about orders placed in a store including products, prices, ship dates, geographical information, and profits. Interesting trends in this dataset have been very well studied, and participants will use SEEDB to quickly re-identify these trends.
- **Election Contribution dataset** [?]: This is an example of a dataset typically analyzed by non-expert data analysts like journalists or historians. With this dataset, we demonstrate how non-experts can use SEEDB to quickly arrive at interesting visualizations.
- **Medical dataset** [?]: This real-world dataset exemplifies a dataset that a clinical researcher might use. The schema of the dataset is significantly complex and it is of larger size.
- **Synthetic data**: We provide a set of synthetic datasets with varying sizes, number of attributes, and data distributions to help participants evaluate SEEDB performance on diverse datasets.

**Scenario 1: Determine Utility.** Participants are provided with three diverse, real-world datasets to explore using SEEDB. For each dataset, participants can issue ad-hoc or pre-formulated queries to SEEDB. SEEDB will then intelligently explore the view space and optimize query execution to return the most interesting visualizations with low latency. Participants can examine the returned queries visually, via the associated view metadata, and via drill-downs. To aid the evaluation of visualizations, the demo system will be configured to also show the user “bad” views (views with low utility) that were not selected by SEEDB. Similarly, we provide pre-selected queries (and previously known information about their results) to allow participants to confirm that SEEDB does indeed reproduce known information about these queries. Participants will also be able to experiment with a variety of distance metrics for computing utility and observe the effects on the resulting views.

**Scenario 2: Determine Performance and Optimizations.** This scenario will use an enhanced user interface and synthetic datasets mentioned above. Participants will be able to easily experiment with a range of synthetic datasets and input queries by adjusting various “knobs” such as data size, number of attributes, and data distribution. In addition, participants will also be able to select the optimizations that SEEDB applies and observe the effect on response times and accuracy.

Thus, through our study of SEEDB we seek to validate the following statements:

- it is possible to automate labor-intensive parts of data analysis
- aggregate and grouping-based views are a powerful means to identify interesting trends in data
- the right set of optimizations can enable real-time data analysis of large datasets

## 8. CONCLUSION

In the paper, we discussed SEEDB : a system that provides analysts with visualizations highlighting interesting aspects of the query result. We defined several concrete problems in SEEDB and their solutions, ranging from multi-query optimization and approx-

imation to multi-criteria optimization.